

Yours Truly

Sunil Template



Contents

1	The Convergence of Enterprise, Internet Scale, and High Performance Computing Storage Infrastructures	1
	<i>Jay Lofstead, Eric Barton, Matthew Curry, Carlos Maltzahn, Robert Ross, and Craig Ulmer</i>	
1.1	Introduction	2
1.2	Existing File Systems	4
1.3	Bridging Technologies	5
1.4	Object-Based Stores	6
1.4.1	HPC Oriented Object Stores	6
1.5	Next Generation HPC Storage Systems	7
1.5.1	Lustre/DAOS	7
1.5.1.1	Containers and Objects	8
1.5.1.2	Transactions	8
1.5.1.3	API Complexity	9
1.5.1.4	Discussion	10
1.5.2	Kelpie/Data Warehousing	10
1.5.2.1	Discussion	12
1.5.3	Hybrid Models	12
1.5.3.1	Discussion	14
1.5.4	Radical Departure	14
1.5.4.1	Logical Structure	15
1.5.4.2	Storage Devices	16
1.5.4.3	Discussion	18
1.6	Conclusions	19
	Bibliography	21



Chapter 1

The Convergence of Enterprise, Internet Scale, and High Performance Computing Storage Infrastructures

Jay Lofstead
Sandia National Laboratories

Eric Barton
Intel

Matthew Curry
Sandia National Laboratories

Carlos Maltzahn
University of California, Santa Cruz

Robert Ross
Argonne National Laboratory

Craig Ulmer
Sandia National Laboratories

Abstract	2
1.1 Introduction	2
1.2 Existing File Systems	4
1.3 Bridging Technologies	5
1.4 Object-Based Stores	5
1.4.1 HPC Oriented Object Stores	6
1.5 Next Generation HPC Storage Systems	7
1.5.1 Lustre/DAOS	7
1.5.1.1 Containers and Objects	8
1.5.1.2 Transactions	8
1.5.1.3 API Complexity	9
1.5.1.4 Discussion	10
1.5.2 Kelpie/Data Warehousing	10

1.5.2.1	Discussion	11
1.5.3	Hybrid Models	12
1.5.3.1	Discussion	13
1.5.4	Radical Departure	14
1.5.4.1	Logical Structure	15
1.5.4.2	Storage Devices	16
1.5.4.3	Discussion	18
1.6	Conclusions	19
	Acknowledgements	19

Abstract

Large scale storage infrastructures have been significantly impacted by the growth in data analytics applications. High Performance Computing storage infrastructures, once the extreme end of the storage scale spectrum, must now adapt to technologies optimized for large scale data analytics applications. Hardware changes, such as storage class memory, are also affecting how the exascale storage stack will be constructed. We examine use cases, trends, convergent technologies, and new opportunities generated by this technology blending.

1.1 Introduction

HPC infrastructures have grown around the requirement to handle large, decomposed data structures for parallel computation. Single data objects may be hundreds of terabytes spread across an entire machine. Parallel storage systems have grown trying to address performance and storage requirements while maintaining backwards compatibility with the standard POSIX interface and semantics. Unfortunately, this is proving increasingly difficult, as the POSIX specification was not designed to efficiently support parallel storage [8].

Big data and internet-scale applications, on the other hand, focus on searching through immense volumes of small, loosely associated items looking for patterns or correlations that may lead to insights. Some science applications, such as genomics, have a workload pattern similar to these big data applications. Parallel file systems are not well suited to these workloads because the broad, relatively continuous read demand of independent items does not benefit from the coordination overheads of parallel file systems. Instead, distributing loosely coordinated storage across the compute infrastructure makes more sense. With pressures to effectively leverage a single platform for these

disparate workloads and the shifting storage market, new considerations for how to design storage resources for extreme scale compute systems must be made.

Traditionally, the HPC market has focused on supporting coherent and consistent output methods from parallel sources to parallel targets. This requirement is driven from validating that the output of a single item is complete and correct. Largely, the workload is write-intensive during the expensive, at scale computation process with a read-intensive phase lasting months on cheaper machines or at small scale with low priority. File systems like the dominant Lustre [4] and GPFS [17] systems have been carefully optimized to address these workloads.

The big data market has opposite priorities. The big computation phase requires reading large quantities of data for processing at scale. The output from this process can be handled at much smaller scale later and is orders of magnitude smaller. Given the read-dominant focus, the overhead inherent in coherent and consistent storage for write intensive workloads is both unnecessary and a heavy cost. Instead, systems like HDFS [19] dominate. These work by storing files that will be read for processing throughout the compute area, including the assumption that storage failures are common, prompting replication.

The output from initial stream or file processing for the big data workloads use distributed object-based storage technology. It offers independent, uncoordinated data access with a simplistic key search space for subsequent analysis. The profit potential for this market has caused an explosion in specialized products aimed at accelerating this processing style. For small enough data sets, in memory object stores like memcached and Radius dominate. For larger data sets, approaches like Google's Big Table are accomplishing the same function. There are also hardware products targeting this market segment, such as Kinetic [18], offering a native object interface for the devices connected directly to a network.

Adding complexity to this storage environment is the relentless performance improvements and cost reductions for solid state storage, like NAND-based flash memory. These devices have already rendered 15,000 RPM disk drives obsolete. The 10,000 RPM disk drives will not survive for more than a few more years. New disk technology like shingled drives [23] offer a path for disks to survive longer. The enormous capacities for read intensive, write infrequently workloads is very attractive for many communities. For example, storing images created sequentially for later read-intensive processing can yield a better cost/performance balance.

This chapter investigates these new technologies and how they affect extreme scale computing. We evaluate how the HPC environment can and must adapt to this new storage environment to address future computing needs and to take advantage of the different kinds of technology being developed. We will also consider the planned reintegration of large scale computing from the split of big data applications from simulation-based computing with both

the necessary and forced integration of these large, expensive platforms for multi-use.

1.2 Existing File Systems

HDFS developed to support the Hadoop implementation of the MapReduce system. It offers a distributed, replicated file store optimized to support the MapReduce processing configuration. Ceph [21] also addresses this distributed computing infrastructure, but with a different emphasis. It seeks to offer scale out performance for objects across a storage infrastructure with unreliable storage. However, Ceph has historically had some fundamental assumptions that incur overheads that are unacceptably large for the high end HPC center [20]: Traditional HPC file systems rely on large RAID arrays for reliability while Ceph's design favors small storage targets with no hardware-supported redundancy but instead imbued with Ceph's peer-to-peer technology to allow self-healing. With the rise of cloud systems using object-based storage, such as Amazon's S3, the interaction style offered by Ceph has been adopted for similar workloads. Ceph's unreliable storage handling works for both new storage additions as well as failures and recoveries without interruption. Ceph offers a complete file system including metadata management support as well as an object system for users. GlusterFS [3] focuses on providing a network attached storage interface to storage distributed throughout a cluster. Rather than providing metadata services, GlusterFS relies on the underlying storage file system for most basic capabilities, such as security.

Parallel file systems are optimized to support large files that must be spread across multiple devices. For example, a 100 TB file cannot fit on any current storage device and cannot be stored with any performance. Parallel file systems solve this problem by using multiple devices spread across multiple servers together as if they are a single device. Data is striped across devices, all of which can be written to or read from simultaneously. This parallel access offers aggregate performance enabling manipulating very large files with reasonable performance. Because of these characteristics, parallel file systems are deployed on most simulation intensive large scale compute systems to handle the large single object output characteristic of these applications. Lustre is arguably the most popular parallel file system appearing on a majority of the Top500 machines. IBM's GPFS is increasing in popularity as optimizations focusing on addressing big data workloads are incorporated. PVFS [16] offers a rethinking of some of Lustre's earlier limitations to give better scalability characteristics.

Discussion

The different optimizations distributed vs. parallel file systems offer are at

the cost of supporting the other kind of workload. As was mentioned above, parallel file systems aim to support very large objects and aggregate simultaneous writing and reading for a single object across the array. For workloads consisting of entirely small objects, this functionality and overhead is a cost. Similarly, the inability of the distributed file systems to handle arbitrarily large objects and massive parallel simultaneous access to a single object make them unsuitable for simulation workloads.

For both workloads, a more flexible object-based interface are being considered. This is discussed more below.

1.3 Bridging Technologies

While the shift to object stores pushes forward, we do have several pieces of bridging technologies that may transition easily into the object storage environment.

For efficient checkpoint/restart storage, two similar approaches are leveraging the complex storage hierarchy to keep this rarely needed data out of the storage array. Scalable Checkpoint Restart (SCR) [14] and the Fault Tolerant Interface for Hybrid Systems (FTI) [2] both use other in compute area resources to address checkpoint writing and recovery performance. SCR focuses on hierarchical modeling while FTI focuses on supporting hybrid compute like GPU-based systems.

Also working to bridge the gap are burst buffers. In short, burst buffers are some generally non-volatile memory (NVM) somewhere within the compute area intended for use as intermediate storage of some sort. One of the first formal burst buffer products is DataWarp from Cray. IBM also has a similar product to be deployed on machines in 2016. The goal of these technologies is to offer a high performance, but limited capacity storage area as a fast cache for the parallel file system. Other uses for this storage is not the primary concern. The motivating use case is that buying sufficient disks to achieve the necessary parallel, aggregate performance is too costly, takes up too much space, and has too many failures just to the very large number of devices. Instead, disk is bought primarily for capacity while the burst buffer is bought based on performance. The combination price yields less total storage, but the performance characteristics are more favorable. Moving to a pure NVM solution directly is too cost prohibitive for the next several years. Capacity and durability are still too expensive.

1.4 Object-Based Stores

The earliest object store is probably the Wisconsin Storage System [5] published in 1985. It offered a general storage infrastructure for both databases and file systems. Many current systems were built using a similar infrastructure, such as Lustre [4], GPFS [17], and Panasas [22]. The general idea is to offer a standardized way to address an arbitrary storage space with a key-based access. These object storage systems assume that some sort of structure is imposed on top to track what objects correspond to which user items.

While object storage may have been used behind the scenes for years, raw object storage exposed to the end-user programmer did not come into vogue until the big data era arrived. System developers for big data processing realized that the overhead imposed by the object management forced serialized or at least coordinated access. By shifting the mapping load to the end-user programmer level and using the object storage layer directly, greater perceived performance can be achieved. In many cases, by stripping down the requirements to the absolute minimum required semantics for a particular application, actual performance gains are achieved. The explosion of specialized storage systems like HDFS and GFS represent this model. Key for this model achieving performance is the ability of each process to work independently without any required consistency or coherence with neighbor processes potentially working on part of the same data set. The dominant object stores are systems like Memcached [6]. This stands in stark contrast to how supercomputing applications generally operate.

The supercomputing domain maintained the consistency requirements due to the bulk synchronous parallel processing. Instead, parallel, that is coordinated, file systems were embraced. The challenge today is that parallel file systems are having difficulty effectively scaling to handle the IO demands.

Here we want to talk about how object-based key-value stores are used for big data applications summarizing the specific features that identify this market segment.

1.4.1 HPC Oriented Object Stores

Parallel file systems inherently have an object-like layer beneath the surface. The requirement to spread a single file across multiple devices for capacity reasons alone prompts this approach. The actual implementation may vary, such as using individual files within a local file system, each representing part of parallel file. Popular examples include Lustre [4] and GPFS [17].

In some cases, directly using a key-value store for HPC applications is being considered [24]. The next generation Lustre project is also considering a key-value infrastructure [1] to address performance challenges.

The real challenge with key-value stores for HPC applications is the meta-

data management. All of these projects have taken a similar step to the big data application is that the applications are required to manage the object list to determine what data is stored in which object.

1.5 Next Generation HPC Storage Systems

Next generation storage infrastructures are taking on several different forms with additional approaches both possible and likely. Almost certainly, multiple of these approaches will last including evolutions of existing systems such as IBM's GPFS. This section discusses four examples of different kinds of storage infrastructures for the exascale environment.

1.5.1 Lustre/DAOS

The US Department of Energy is using a "Fast Forward" series of projects to kickstart needed developments in the vendor space. One such project is the Storage and I/O (FFSIO). The basis for this is extending the existing Lustre infrastructure to offer an object interface, incorporate support for burst buffer-like technologies, and demonstrate use through a client I/O library. The first phase completed in 2014 and the second phase started in late 2015.

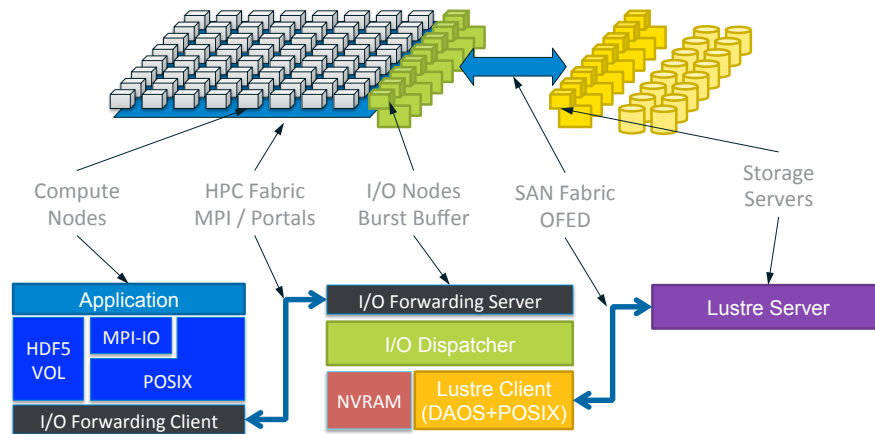


FIGURE 1.1: FFSIO Architecture and Component Mapping

Overall, the system design (see Figure 1.1) consists of a friendly end-user API that manipulates either the I/O Dispatcher layer or the DAOS (Distributed Asynchronous Object Storage) layer.

The first phase design focuses on a few themes. First, a container and object structure is assumed. Second, transactions are integral to offering de-

coupled performance. And third, hide additional API complexity behind a friendly end-user API.

1.5.1.1 Containers and Objects

FFSIO presents a object-like interface that is mostly backwards compatible with the standard POSIX files that an end-user API like HDF5 would create. It has containers that represent files and objects of different types to represent the different kinds of data stored for a “file”. To keep a pure view on the system, there is no master container or any other such construct. Instead, it is assumed that a user or system will specify a standard well-known container that contains a list of available other containers. Some additional complexity is introduced through this approach, but the decoupled requirement to update a central metadata service for any file creation or access operation is eliminated. Instead, either the user or policy would have to keep track of available containers. The reasoning behind this structure is a bit complicated, but quite reasonable.

Potentially extreme overheads for writing logically contiguously formatted data is well documented. For example, ParColl [25] showed as much as 90% overhead on as few as 512 processes. For the Six Degrees of Scientific Data [13] work, the authors were unable to perform some performance comparisons because a single data output could not be written in the maximum normal job time. These sorts of impacts are prompting rethinking the logically contiguous format favored by popular APIs like HDF5 and PnetCDF. The perceived advantages when these formats were selected have not proven to scale to very large, 3-D domain decompositions in particular. DAOS, offers the object storage infrastructure while the HDF5 API at the end-user level is maintained almost exactly. The only end-user API changes are related to incorporating transactions. This shift almost exactly maintains backwards compatibility while taking advantage of more scalable infrastructures.

1.5.1.2 Transactions

The system jitter and other causes for processes to fade out of sync affect overall I/O performance as well. While transactions are typically seen as a synchronous operation that blocks while all participants manage state, FFSIO takes a passive approach and manages transactions asynchronously. The key idea is to keep writes in a log-style copy-on-write structure and annotate each write with which transaction it is part of. An overall transaction mapping knows how many writes are expected and can decide when a transaction is complete or not. By monitoring the system status, it can passively detect many failures abort the affected transactions. The potential overheads to this approach are refining the idea, but the core concept shall remain the same.

Transactions are manifest in two different ways. First, the end-user API layer can opt to have the FFSIO system manage the transactions by that layer tracking the status of all clients directly. Second, transactions can be

managed by the end-user layer themselves with a single process indicating the transaction state to the FFSIO layer. The idea for the former approach is that it offers transactional functionality without undue end user code changes. At the IOD and DAOS layer, transactions are handled slightly differently.

The IOD layer is intended as a staging area for data that may or may not need to be written to persistent storage. This is one implementation of the burst buffer concept. The key idea is that the transactions in this layer may or may not ever be pushed down to the DAOS layer. Instead, once the data set is processed, only the analyzed results will be saved eliminating the need to write large, temporary data sets to the slower, but much larger capacity storage array. However, these transactions are connected to those at the DAOS layer.

The DAOS layer is intended to function as a general object store that would replace an existing data center wide shared parallel file system. Anything written to this layer is expected to have longer lifetimes and potentially be accessible from other platforms. Since the transactions work a bit differently, in particular some sequential transactions may be missing, the name is changed to epochs instead. The slightly different idea for an epoch is to represent some important saved state. Since it is likely a revision of an existing version, differencing mechanisms are tested to determine how to efficiently store both versions with less space and minimal reconstruction overhead. When small changes are persisted, the space savings can be enormous.

The Doubly Distributed Transactions Project (D²T) offers a more general approach to transactions [12, 11]. The end-user programming cost is higher to support the more general use. Other than the synchronous requirement for failure detection, the performance is excellent.

For future transactional use in the storage hierarchy, some mixture of the three approaches is likely. The passive transactions will be used for backwards compatibility. The end-user, single process managed transactions will work well for storage system interactions. The D²T approach could be used in isolation or to support the end-user, single process managed transactions.

1.5.1.3 API Complexity

One of the complexities of the FFSIO stack is addressing various scale platforms. For the largest platforms, including an IOD layer makes sense. The particular architecture may vary, but that kind of technology will be important to reduce compute throughput due to waiting for I/O to complete. For smaller platforms, either the IOD layer is deemed unnecessary or it may be an extra expense that offers little advantage. In this case, omitting this component both simplifies the system and offers more direct performance.

The challenge with this approach is that the end-user level must address both the IOD and DAOS APIs directly and be able to choose for fully portable code. The advantages of being able to effectively scale down were maintained by using a higher-level end-user API, like HDF5. For such a professionally

produced library, offering a second version is not a major undertaking. Further, it enhances the HDF5 value by offering greater platform scalability.

1.5.1.4 Discussion

Overall, the shift from a traditional parallel file system to something like the FFSIO system will be a big step forward towards supporting more big data style hardware directly. While the design is not perfect, the second phase is addressing most, if not all, of the detected shortcomings.

1.5.2 Kelpie/Data Warehousing

Kelpie is a distributed, in-memory object store from Sandia National Laboratories that serves as a building block in high-performance computing for implementing custom, data-management services. The fundamental goal of Kelpie is to provide a way for users to decompose their complex datasets into data objects that the library can move between nodes in a safe but efficient manner. Kelpie provides simple abstractions for dealing with distributed data, and utilizes the Nessie communication library to orchestrate how data migrates between nodes. Nessie provides (1) a low-level RDMA layer that has been ported to different HPC fabrics and (2) an RPC layer that enables users to invoke function calls and initiate RDMA transfers on remote nodes. Kelpie uses the former to make an application's data objects accessible by the network interface, and the latter to coordinate data handoffs between nodes.

Kelpie manages data objects for applications, where a data object is a simply a contiguous block of application data that is labeled with a globally-unique key. The contiguous constraint is necessary because Kelpie registers the memory with the network interface, which in turn allows the hardware to RDMA the data without involving the kernel in virtual to physical address translation. The key used to label an object has three components: an application-specific identifier and a two-dimensional user label. The application-specific identifier enables higher-level services to house different datasets in Kelpie with isolation. Users are not required to use the second dimension of the user label portion of the key. However, the second dimension is often useful for grouping related items together in the store. For example when storing complex mesh datasets in Kelpie, the first dimension of the key (or row) may be used to describe a particular region of a mesh. The second dimension of the key can therefore be used to organize different variables associated with the region (e.g., node locations, pressure, or temperature). This approach allows each variable's data to be stored in its own, independent block of contiguous memory, and provides an opportunity for a user to easily downselect the items they need when retrieving data from Kelpie.

Kelpie nodes are equipped with a Local Key/Value (LKV) structure for managing data objects that are available in the local node. The LKV performs three important tasks. First the LKV provides a means of tracking data ob-

jects that are currently in transit and protect the system from deallocating an object before remote nodes have finished transferring it. Second, the LKV structure provides a means for data to be staged at a remote node with only trivial involvement from the destination. For example, an application may push multiple data objects to a node in anticipation of work that will be scheduled on the remote node at a future time. Finally, the LKV structure provides a location for applications to store callbacks to execute when data does arrive at a node. These callbacks make the data store more active and are fundamental to applications that are highly asynchronous or event driven.

In order to address scalability concerns, data management services for HPC typically decompose their work in a hierarchical manner that maps ownership of different portions of the dataset to different groups of nodes that are close in proximity. Kelpie provides the ability to assemble multiple teams of nodes together through a resource management interface. This interface allows users to create and reference a team of nodes together to function as a single data resource that employs a standard data API. A resource has three components: a path-like name that allows a resource to be globally referenced and located by the runtime, a list of physical nodes that implement the resource, and client-side software that defines that policy for how data is managed in the resource. Kelpie provides common implementations of resources, but is easily extended with user-defined modules. A distributed hash table (DHT) is an example of a commonly-used resource in Kelpie. A DHT is composed of N Kelpie servers where data is distributed by using a hash of the first dimension of the key to specify which node should store the object. The DHT resource client software simply maintains the list of servers and then references the proper destination when a user performs put or get operations. Users with reliability concerns have made similar resources just by creating new client software that stores a data object to the hashed server, and a replicated copy to its neighbor. While resources may contain API extensions to support higher levels of functionality, they all support basic communication primitives that enable users to swap one type of resource in for another in most cases.

An example of how Kelpie is being utilized to support data management for higher-level applications can be found in Sandia's DARMA project. DARMA is developing an asynchronous, many-task (AMT) approach to computing that will help codes scale to next-generation platforms. AMT codes specify their execution in the form of a large, directed acyclic graph of tasks (or task DAG) that can be scheduled by a runtime on distributed resources. The data objects consumed and produced by each task form a dependency graph that dictate when and where work can be scheduled. By utilizing Kelpie as the mechanism by which data objects are migrated between different nodes and resources in the system, DARMA can focus on the complex job of making policy decisions about how the work should be orchestrated in the platform.

1.5.2.1 Discussion

Kelpie is clearly taking a different approach from the FFSIO project. By pushing the key-value structure all the way up to the application, it requires a whole new interaction between the applications and the storage infrastructure, but at a potential advantage of much richer, higher performance. For example, by integrating with the AMT management layer, Kelpie can predictively move data as necessary or even guide the AMT management layer to place computation in a different location to avoid the data movement. Also, by offering object level access beyond what an application would normally write to storage, potentially new data access routines for richer analysis could be integrated without requiring any application changes. Simply by registering all data through Kelpie for use in the AMT management layer, other users could take advantage of the metadata and select data for additional analysis without worrying about explicit staging or writing to some persistent store.

1.5.3 Hybrid Models

A different take on the problem is being investigated in a DOE sponsored project called SIRIUS primarily by Oak Ridge, Sandia, UC Santa Cruz, Rutgers, and Brown. In this project, we seek to address the performance mismatch more directly through several approaches simultaneously.

First, using a single storage tier at a time leaves potential performance untapped. Instead of the traditional structure where an output is pushed to a particular storage tier and it is maintained on that tier (or partially through some sort of hierarchical storage management system), all tiers are used directly at the same time. For example, when writing a large data set for a molecular dynamics code simulating crack formation, only certain areas contain important features required for detailed analysis the rest of the data is not as important. Consequently, SIRIUS seeks to write the important parts of the data set to the fastest tier with the less important spooled to slower tier(s). In many cases, if the output for the slower tier fails, the data set is still usable enabling throwing away the less interesting data set portions given system performance pressures.

For performance reasons, much like the other solutions discussed above, SIRIUS is assuming an object storage layer beneath the API. This layer could be DAOS, Kelpie, the Sandia Sirocco storage system, Ceph, or any other object-based infrastructures. Which object storage infrastructure employed will be flexible depending on the data center and application workloads requirements.

Second, feature identification and extraction can be expensive in time and space or may not be applicable to the simulation type. For example, particle in cell simulations use statistical measures for all particles in cell regions to determine simulation values. Tagging particular particles as potentially extraneous for a particular output is expensive. Instead, other techniques to slice

the data may be employed. For example, simple data rearrangement may yield sufficient quality for most purposes dramatically reducing the data volumes used in visualizations by a factor of more than 100x. By selecting every fourth element in a mesh and then only selecting the 3 bytes required to represent the sign, exponent, and the most significant bits of the mantissa is a radical data reduction. This small data set version could be written to fast storage while the other possible slicings can be written to slower storage tiers that presumably have more capacity. By being able to slice data for different placement and offering the significant data reductions with little loss in analysis quality, time to insight can be improved. Should the full quality data set ever be required, all of the pieces can be reassembled from the various storage tiers and processed as a single unit. Alternatively, an auditor offers a lossy approach.

Auditors are an idea that recognizes that while simulations have refined over the years, much faster, simpler, older models can still be used over short timeframes to get a very close approximation to the fine grained, fully featured simulation. Using these older models as an auditor for the full fidelity simulation, we can track, only over relatively short time and space constraints, with radically less computational load and data sizes. By using an auditor running along side the full fidelity simulation, the error magnitude can be managed resetting the auditor when the drift is unacceptably large, approximate data quality can be maintained. Using statistical techniques, auditor output can be used to regenerate statistically similar full fidelity simulation values within a known error bounds forced by the state comparison to determine the auditor drift. Using this approach major events or periodic timesteps can be output using full fidelity while many intermediate results can be output using the auditor state generating high quality lossy representations of what the intermediate simulation state would be like. There are other lossy approaches available. For example, ISABELA [10] offers a sorted data set with B-Spline curve fitting to reduce the data set size. Delta [7] looks at simple diffs to reduce the data set with low space-time overheads with no loss. Wavelet compression [9] low space-time overheads with potentially radially high data compression with predictable error rates. Each of these approaches can be useful in different circumstances.

Third, SIRIUS focuses on Quality of Service (QoS) to help guide users to choosing an acceptable time/quality tradeoff for a given operation. For example, in the data reorganization example above, using the every fourth point, 3 bytes of 8 bytes values may be able to be loaded and analyzed in 1 minutes. Reading the full fidelity may take 2 hours. For a quick peak to see if the data might be “interesting”, the 1 minute overhead is likely enough. Should a general feature seem to be present, a data set could be noted for later fully detailed review. The 2 hour overhead could be scheduled rather than delaying the interactive scientist.

1.5.3.1 Discussion

SIRIUS seeks to take advantage of the whole storage hierarchy to get maximum I/O bandwidth. By incorporating techniques to slice, compress, or lossy compress the data, more efficient storage tier use can be achieved. Since only the small, most important part of a data set is stored in the fast, small, expensive tier, many more outputs or different applications can use the tier without strong considerations about exhausting the capacity. By including QoS measures as part of the system enabling users to interact with requests to determine the quality/time tradeoff, end user productivity can be enhanced with little to no loss of scientific validity. For final report and paper submissions, the time can be spent for full fidelity data reads and analysis to validate the results generated from the lower fidelity data versions. These kinds of techniques will be increasingly important to maximize both platform productivity as well as scientist productivity.

1.5.4 Radical Departure

The Sirocco Storage System from Sandia takes inspiration from peer-to-peer systems to rethink how a storage system should work for HPC. Instead of relying on existing notions of striping files across storage devices, it is fully object based and works under the idea of primarily caching under data resiliency requirements.

Existing parallel file systems, such as Lustre, GPFS, Panasas, and PVFS/OrangeFS, all offer a way to store objects larger than a single storage device by striping data across devices simultaneously enhancing performance. Each of these systems have optimized in a different way. Lustre exposes all of the tuning parameters allowing users to tweak settings to optimize performance. The downside is that these parameters must be managed to achieve good performance. GPFS hides most of these parameters and attempts to manage performance optimizations with moderate success. Panasas dynamically increases the number of storage targets used for a file as the simultaneous parallel writer count increases. PVFS/OrangeFS has optimized metadata operations reducing the load on the metadata server.

Sirocco seeks to completely rethink this model to move away from the rigid striping model with a separate metadata service into a storage fabric with an inherent assumption that resources may be transient. This fundamental rethinking is being done with a nod to backwards compatibility by offering a POSIX interface that can access a particular view of the storage system. The native interface is object based with containers collecting objects each of which have multiple data forks that can be accessed through an address space. These four abstract levels offers flexibility to address HPC, Experimental and Observational Data, and large scale data analytics needs.

The roots of Sirocco are in the Lightweight File Systems (LWFS) [15] project at Sandia from several years ago. LWFS sought to strip down a file

system to the bare, required components and allow users to add additional capabilities paying the overhead for only those features required. For example, the LWFS core consists of an object store with authentication and authorization services only. Other features, such as naming and consistency control, are left to separate services. Sirocco seeks to provide the storage layer that an LWFS-style system would require. Keeping this in mind will help explain some of the decisions made. For example, a non-Sirocco client accessing or manipulating data is equivalent to going directly to disk and reading or modifying bytes rather than going through the file system interface. Protections against such access are not included since it is intended to act like a device like disk only accessed through a file system interface of some sort.

One way to think about how Sirocco works is as a reverse Bittorrent system. In a typical Bittorrent system, a seed offers a file for others to replicate as they request it. A tracker offers a directory to find where seeds are located so that new clients can discover what is currently available. Over time, a requested seed will be replicated to the local client based on pulling pieces from a variety of sources currently offering a copy of the file. Which seeds are available over time change dynamically making it likely that a new copy of a file came from multiple sources, some of which were only available for part of the transfer time.

Sirocco takes a reverse approach for writing because a client wants to store a file into the server collection with some resilience properties. As it is pushed into the space, it will be copied and replicated according to the device characteristics and availability and resilience requirements asserted as the data is pushed.

For reading operation, Sirocco will go find the current data version and return it to the requesting user. For performance reasons, a user could force a copy to migrate to a particular location. The old version may be flushed on a space required and resilience considering basis. The collection of Sirocco servers is referred to as a swarm further reinforcing the Bittorrent analog.

To understand how Sirocco works at a detailed level, several concepts and how they work in Sirocco are required. Each of these are explored below.

1.5.4.1 Logical Structure

Logically, Sirocco uses a containers, objects, forks, and address spaces tuple to identify a particular byte. Each of these concepts plays a different role offering specific features.

All of these concepts rely on IDs or ranges. Each of these are represented by a 64-bit number. This specifies the size of all of these spaces. When considering these concepts, a rich format like HDF5 helps understand the purpose of each. Sirocco goes beyond HDF5's capabilities to offer new features common in other kinds of storage systems.

Containers

The container is equivalent to a file in a traditional file system. The idea of using a container rather than a file is that it is really a capability envelope for an object collection rather than a single byte stream. Each container contains zero or more objects.

Objects

Each object can be thought of as a variable in an HDF5 file. For example, an array or scalar value may each be stored in a single object. The object 0 for every container is reserved.

Forks and Address Spaces

Each object can have multiple versions or views referred to as a “fork”. These forks are similar to the old Macintosh OS resource and data forks stored for each file. In this case, the number of forks is limited by a 64-bit number. The fork 0 for every object is reserved for information for the security system. This highlights that forks can serve multiple purposes beyond simply being a resource and a data fork for a file. In this case, there is a security information fork. Another use would be to specify a compression level for a version of data. If a particular compression type or level is available, it is stored in a well known fork ID. This can also be used for versioning to offer a rudimentary source control system or represent the old VAX VMS file versions concept. Many other potential uses are also possible and not artificially limited by the system.

To access data within a fork, a 64-bit address space is referenced. Each selected address portion is called a *range*. When writing, data is provided for the range and pushed into Sirocco. The address space is treated as a key value store using the address ranges as the key and the data as the value. Each value is versioned with the versions being accessible.

Each address space within a fork is limited to 4 GiB currently. For the short term, users are required to write multiple ranges as part of an extent to write a larger space.

1.5.4.2 Storage Devices

Storage devices are represented by a Sirocco interface. Each device has its own Sirocco interface and has particular characteristics that determine how it is used. For example, a RAM backing store may be used for a compute-node Sirocco server or a disk may back a different Sirocco server. How and when these are used is driven from the resilience characteristics requested for an extent. Device collections that share the same characteristics can be thought of as comprising a single storage tier. Sirocco makes no such association as it treats all potential storage locations as peers and selects which to use based on the resilience characteristics requested.

Device Characteristics

Device characteristics affect two different system functions. First, they specify the performance that can be expected when accessing the device. Second, they specify a resilience metric that is used for data protection and availability. Both of these work together to guide data placement and migration.

The four basic characteristics are latency, bandwidth, load, and volatility. The first two are exactly as expected for a storage system. Load is a dynamic measure of how busy a storage device is currently. Migrating data to a heavily loaded device is a poor performance choice if other options are available and Sirocco considers this to try to offer the best overall performance. Volatility is a measure of how safe data stored in this device is compared to other devices. RAM vs. node local SSD vs. centralized SSD vs. disk vs. tape all offer different volatility metrics. This metric in particular is used to implement the resilience features.

When an extent is provided to the Sirocco API, two resilience characteristics are provided—short term and long term. The short term requirement describes how to store the data before the Sirocco API returns control to the user. For example, if a two-copy on persistent store short term requirement is specified, Sirocco will ensure that two storage devices meeting this requirement have successfully acknowledged storing the data before returning control to the user.

The long term resilience requirement comes into play as the amount of space available is insufficient requiring data ejection or deletion. Any extent selected for eviction or deletion will be reconciled according to the long term resilience requirements. If the long term requirements specify resilience characteristics typical of a tape system and insufficient additional replicas to meet the resilience requirement exist, Sirocco will migrate the affected data towards or into the tape system to free space required for the current operation.

Finding Data

Given the assumed transient device existence, Sirocco works to keep data with required resilience characteristics as devices come and go. When a device is discovered to be missing, any replicas stored on that device must find alternative locations to preserve the required resilience. News of the missing device is propagated to all participating storage servers so they can react to the loss appropriately. When a new server comes online, a similar propagated message informs the swarm of the new potential target and the characteristics it offers.

Should all replicas for a file go offline effectively simultaneously, it is possible for data to be stored securely, but not be accessible. Sirocco works to minimize this possibility, but it is impossible to guarantee that data will always be available unless all data is replicated on all storage devices—a clearly infeasible solution.

Authoritative Copy

Something that may come to mind with the broad discussion of having multiple data copies for resilience and transient storage devices is how to know

what the “official” data version is. This is handled through the idea of an authoritative data copy. By default, Sirocco selects one copy as the authoritative copy. Should that copy be “far” from the client, the client can request that the authoritative copy migrate to a nearby server to offer better interactive performance.

Server to Server Communication

Since Sirocco has no central point of control, any data can be requested or written to any server in the swarm. If the data is not stored on that server, Sirocco starts to search by asking nearby servers if they have or know where the data is stored. As that message propagates out, the user will eventually be told the data cannot be found or have the data returned to them. For a basic request, the data may or may not migrate a copy to the server the client initiated the request through. By requesting the authoritative copy migrate, the user can force the primary copy to move to the server the request was initiated through.

Client APIs

While the native interaction for the 4-tuple (container ID, object ID, fork ID, address space) is an object interface, it does not address needs of existing applications or file system APIs. While this is the preferred interface since it offers the richest interaction and control mechanisms, backwards compatibility is required. A full POSIX interface has been implemented, but it is limited. For example, a separate metadata service will be required for naming that maps to containers. Also with this mapping, how to handle objects, forks, addresses, and versions is required. By default, it serializes the objects and forks, but only the most recent versions.

1.5.4.3 Discussion

The four simple 64-bit identifiers offers tremendous capacity and storage system flexibility that should be able to address high end storage needs for years to come. With these fixed numeric spaces, but no incorporated metadata system, some operations are no longer apparent. For example, since all IDs exist in theory and storage resources may come and go, the idea of “creating” or “deleting” something is alien. When a particular tuple is specified, as long as the requester has authority to access that space, they will be given access. When reading, if the requested space does not exist in the current space, a “not found” is returned. To delete an item, it must be cleared from the system by doing a write setting the record length to zero.

Since performing since address space writes at a time may be highly inefficient, the concept of a record and an extent exist. A record is the version and length of a value that represents an address range within a fork. An extent is a collection of records to be operated on. Most Sirocco APIs use an extent rather than a record for data manipulation. Processing of extents is in an

undefined order and all have the same version. The extent processing is done as a single transaction.

Sirocco fundamentally rethinks how to offer a storage system for a high end computing environment. By rejecting the current zebra model for a peer-to-peer style model, resilience features can be incorporated more easily. While some of the limitations of this approach may complicate the file system interaction built on top of Sirocco, the flexibility and features make considerations worth the trouble.

1.6 Conclusions

The growth in data analytics and the corresponding embracing of object-based and key-value technologies is dramatically shifting the storage landscape. With pressures to support both traditional HPC simulation workloads and large scale data analytics on the same platforms, the pressures on HPC platforms, software infrastructure, and applications continues to increase. Through adapting programming models, such as using asynchronous many task with a system like Kelpie to support data management, the changed storage features can be fully exposed and embraced by the applications.

For existing applications that either cannot or will not shift to an AMT model, systems like the FFSIO and SIRIUS offer alternatives that with only minimal modification, the applications can achieve many benefits.

Based on our existing efforts and thoughts, the move to the distributed objects and key-value stores will add complexity and incur a loss from parallel file system functionality. However, these costs are proving acceptable given the performance and functionality gains being achieved. As the market and platforms continue to shift, we see HPC applications being better able to take advantage of data analytics infrastructure while the data analytics software will need to make some adjustments to the different platform requirements for HPC simulation-oriented machines. Eventually, both should be able to co-exist with minimal performance loss. Ideally, both will benefit from the combination of the best of both worlds.

Acknowledgements

Sandia National Laboratories is a multi-program laboratory managed and operated by Sandia Corporation, a wholly owned subsidiary of Lockheed Mar-

tin Corporation, for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-AC04-94AL85000.

Bibliography

- [1] Eric Barton. Lustre*-fast forward to exascale. *Lustre User Group Summit*, 2013.
- [2] Leonardo Bautista-Gomez, Seiji Tsuboi, Dimitri Komatitsch, Franck Cappello, Naoya Maruyama, and Satoshi Matsuoka. Fti: high performance fault tolerance interface for hybrid systems. In *Proceedings of 2011 international conference for high performance computing, networking, storage and analysis*, page 32. ACM, 2011.
- [3] Eric B Boyer, Matthew C Broomfield, and Terrell A Perrotti. Glusterfs one storage server to rule them all. Technical report, Los Alamos National Laboratory (LANL), 2012.
- [4] Peter J. Braam. The lustre storage architecture. Cluster File Systems Inc. Architecture, design, and manual for Lustre, November 2002. <http://www.lustre.org/docs/lustre.pdf>.
- [5] H-T. Chou, David J. Dewitt, Randy H. Katz, and Anthony C. Klug. Design and implementation of the wisconsin storage system. *Software: Practice and Experience*, 15(10):943–962, 1985.
- [6] Brad Fitzpatrick. Distributed caching with memcached. *Linux journal*, 2004(124):5, 2004.
- [7] Gregory Jean-Baptiste, Jay Lofstead, and Ron Oldfield. Delta: Data reduction for integrated application workflows. Technical Report SAND2015-5029, Sandia National Laboratories, June 2015.
- [8] Dries Kimpe and Robert Ross. Storage models: Past, present, and future. In Quincey Koziol and Prabhat, editors, *High Performance Parallel I/O*, chapter 30, pages 335–345. Chapman & Hall/CRC, 2014.
- [9] Andreas Klappenecker and Frank U. May. Evolving better wavelet compression schemes. In *in Proc. of Wavelet Applications in Signal and Image Processing III*, 1214, pages 614–622, 1995.
- [10] Sriram Lakshminarasimhan, Neil Shah, Stephane Ethier, Scott Klasky, Rob Latham, Rob Ross, and Nagiza F Samatova. Compressing the incompressible with isabela: In-situ reduction of spatio-temporal data. In *Euro-Par 2011 Parallel Processing*, pages 366–379. Springer, 2011.

- [11] J. Lofstead, J. Dayal, I. Jimenez, and C. Maltzahn. Efficient, failure resilient transactions for parallel and distributed computing. In *Data Intensive Scalable Computing Systems (DISCS), 2014 International Workshop on*, pages 17–24, Nov 2014.
- [12] Jay Lofstead, Jai Dayal, Karsten Schwan, and Ron Oldfield. D2t: Doubly distributed transactions for high performance and distributed computing. In *IEEE Cluster Conference*, Beijing, China, September 2012.
- [13] Jay Lofstead, Milo Polte, Garth Gibson, Scott Klasky, Karsten Schwan, Ron Oldfield, Matthew Wolf, and Qing Liu. Six degrees of scientific data: Reading patterns for extreme scale science io. In *In Proceedings of High Performance and Distributed Computing*, 2011.
- [14] Adam Moody, Greg Bronevetsky, Kathryn Mohror, and Bronis R De Supinski. Design, modeling, and evaluation of a scalable multi-level checkpointing system. In *High Performance Computing, Networking, Storage and Analysis (SC), 2010 International Conference for*, pages 1–11. IEEE, 2010.
- [15] Ron A. Oldfield, Patrick Widener, Arthur B. Maccabe, Lee Ward, and Todd Kordenbrock. Efficient data-movement for lightweight I/O. Barcelona, Spain, September 2006.
- [16] Robert Ross and Robert Latham. Pvfs: A parallel file system. In *Proceedings of the 2006 ACM/IEEE Conference on Supercomputing*, SC '06, New York, NY, USA, 2006. ACM.
- [17] Frank Schmuck and Roger Haskin. GPFS: A shared-disk file system for large computing clusters. In *Proceedings of the USENIX FAST '02 Conference on File and Storage Technologies*, pages 231–244, Monterey, CA, January 2002. USENIX Association.
- [18] Seagate. The seagate kinetic open storage vision. <http://www.seagate.com/tech-insights/kinetic-vision-how-seagate-new-developer-tools-meets-the-needs-of-cloud-storage-platforms-master-ti/>, 2014.
- [19] Konstantin Shvachko, Hairong Kuang, Sanjay Radia, and Robert Chansler. The hadoop distributed file system. In *Proceedings of the 2010 IEEE 26th Symposium on Mass Storage Systems and Technologies (MSST)*, MSST '10, pages 1–10, Washington, DC, USA, 2010. IEEE Computer Society.
- [20] Feiyi Wang, Mark Nelson, Sarp Oral, Scott Atchley, Sage Weil, Bradley W Settlemyer, Blake Caldwell, and Jason Hill. Performance and scalability evaluation of the ceph parallel file system. In *Proceedings of the 8th Parallel Data Storage Workshop*, pages 14–19. ACM, 2013.

- [21] Sage A. Weil, Scott A. Brandt, Ethan L. Miller, Darrell D. E. Long, and Carlos Maltzahn. Ceph: A scalable, high-performance distributed file system. In *Proceedings of the 7th Symposium on Operating Systems Design and Implementation*, OSDI '06, pages 307–320, Berkeley, CA, USA, 2006. USENIX Association.
- [22] Brent Welch, Marc Unangst, Zainul Abbasi, Garth A. Gibson, Brian Mueller, Jason Small, Jim Zelenka, and Bin Zhou. Scalable performance of the panasas parallel file system. In Mary Baker and Erik Riedel, editors, *Proceedings of the USENIX FAST'08 Conference on File and Storage Technologies*, pages 17–33. USENIX, February 2008.
- [23] Roger Wood, Mason Williams, Aleksandar Kavcic, and Jim Miles. The feasibility of magnetic recording at 10 terabits per square inch on conventional media. *Magnetics, IEEE Transactions on*, 45(2):917–923, 2009.
- [24] Yanlong Yin, Antonios Kougkas, Kun Feng, Hassan Eslami, Yin Lu, Xian-He Sun, Rajeev Thakur, and William Gropp. Rethinking key-value store for parallel i/o optimization. In *Data Intensive Scalable Computing Systems (DISCS), 2014 International Workshop on*, pages 33–40. IEEE, 2014.
- [25] Weikuan Yu and Jeffrey Vetter. ParColl: Partitioned collective I/O on the cray XT. *Parallel Processing, International Conference on*, 0:562–569, 2008.